

# Hidden Markov Models

Jim Glass / MIT 6.806-6.864 / Spring 2021

1

## Sequential Labeling

---

- Human language is fundamentally sequential in nature
- Many NLP tasks involve converting one sequence into another:
  - Part-of-speech tagging
  - Named entity recognition
  - Machine translation
  - Speech recognition
- A range of ML techniques apply to sequence-to-sequence tasks:
  - Hidden Markov models
  - Conditional random fields
  - Recurrent neural networks

2

## Today's HMM Storyline


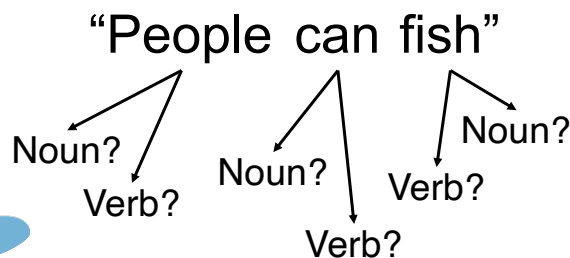
- Part-of-speech tagging
  - Dynamic programming
  - Viterbi search
- Hidden Markov models
  1. Scoring: Forward-backward algorithm
  2. Matching: Viterbi search
  3. Training: Baum-Welch parameter estimation



3

## Part-of-Speech Tagging

- POS tagging assigns each word in a sentence a grammatical tag
  - The tag depends on the word and its context (e.g., sentence)
  - POS inventory is language and corpus dependent
  - Typically used for features, or a precursor for other tasks (e.g., parsing)
  - POS tagging also known as word category disambiguation
- An inherent challenge for POS tagging is word category ambiguity

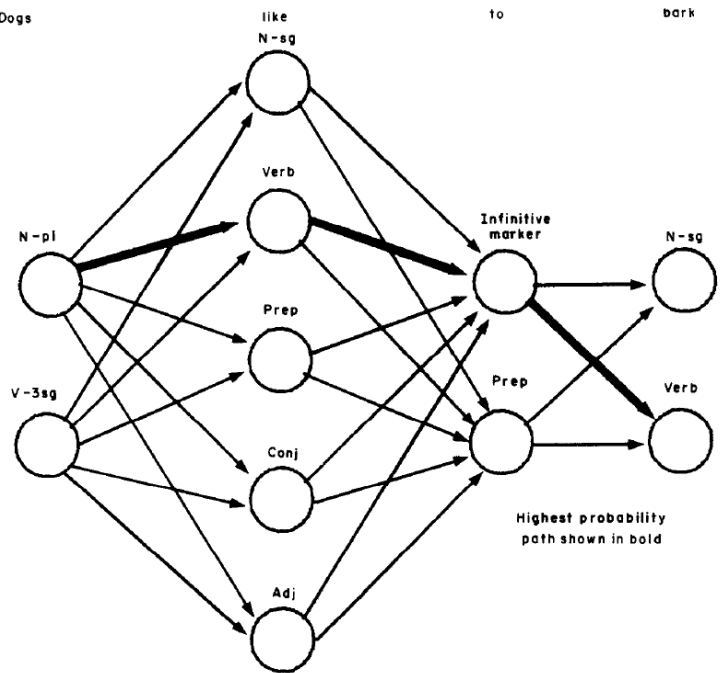


4

## Robust part-of-speech tagging using a hidden Markov model

**Julian Kupiec**

Xerox Palo Alto Research Center, 3333 Coyote Hill Road, Palo U.S.A.



5

## A Probabilistic Formulation for POS Tagging

- Define words  $W = \{w_1, \dots, w_n\}$  and corresponding tags  $T = \{t_1, \dots, t_n\}$
- Given a word sequence, we infer the “hidden” tag sequence  $T^*$

$$T^* = \arg \max_T P(W, T) \quad \text{where } P(W, T) = P\{w_1, \dots, w_n, t_1, \dots, t_n\}$$

- Using the chain rule, we can rewrite  $P(W, T)$  as

$$P(W, T) = \prod_{i=1}^n P(w_i, t_i | w_1, \dots, w_{i-1}, t_1, \dots, t_{i-1})$$

- By making conditional independence assumptions that  $t_i$  depends only on  $t_{i-1}$ , and  $w_i$  depends only on  $t_i$  we can rewrite  $P(W, T)$  as

$$P(W, T) = \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$$

Observation probabilities  
Transition probabilities

6


## Parameter Estimation and Search

- Observation and transition probabilities can be estimated from annotated data or learned via EM algorithm from unannotated data

		people	can	fish	<>
<b>Observations</b>	<>	$-\infty$	$-\infty$	$-\infty$	0
$\propto \log P(w_i   t_i)$	N	-1	-3	-2	$-\infty$
	V	-5	-1	-4	$-\infty$

		N	V	<>
<b>Transitions</b>	<>	-1	-2	$-\infty$
$\propto \log P(t_i   t_{i-1})$	N	-4	-1	-2
	V	-2	-3	-1

- Search space can be represented as directed acyclic graph (DAG)




Weighted finite-state transducers are effective representations for DAGs

7

## Dynamic Programming (DP)

- DP algorithms such as Viterbi search leverage optimal substructure
  - Let  $\phi(i, j)$  be the best path between nodes  $i$  and  $j$
  - If  $k$  is a node in  $\phi(i, j)$ :  $\phi(i, j) = \{\phi(i, k), \phi(k, j)\}$
  - Let  $\varphi(i, j)$  be the cost of  $\phi(i, j)$  (e.g.,  $-\log \text{prob}$ )

$$\varphi(i, j) = \min_k (\varphi(i, k) + \varphi(k, j))$$

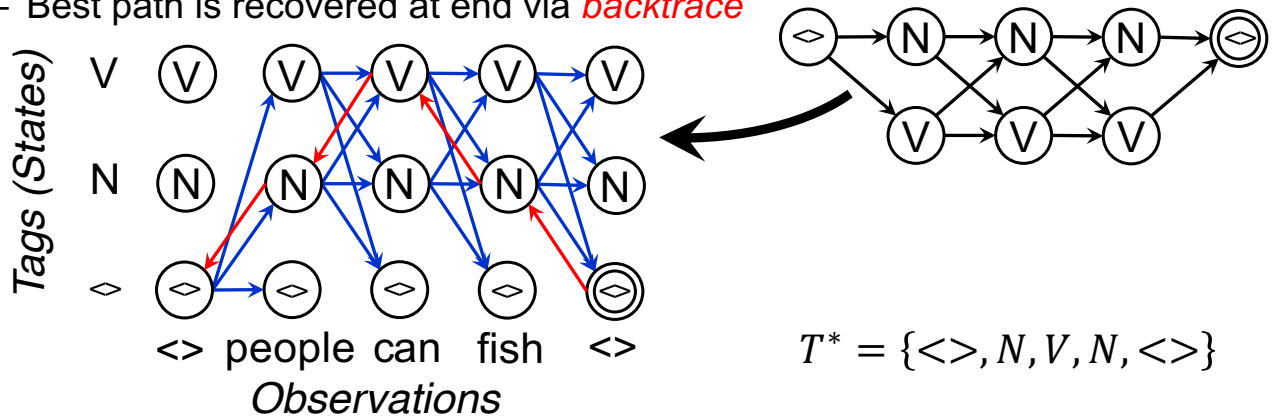


- Solutions to sub-problems need only be computed once
  - Sub-optimal partial paths discarded while staying *admissible*
- Can be implemented efficiently:
  - Node  $k$  retains only best path cost of all  $\varphi(i, k)$
  - Previous best node index needed to recover best path
- Best-first and A\* graph search also leverage optimal substructure

8

## Viterbi Search



- Viterbi search arranges search through a fixed-dimension *trellis*
  - Search advances time-synchronously
  - All partial paths ending at a common node converge at the same moment
  - Per DP, each node retains best score and back pointer to best partial path
  - Best path is recovered at end via *backtrace*



9

## An Introduction to Hidden Markov Models

L. R. Rabiner  
B. H. Juang




IEEE ASSP MAGAZINE JANUARY 1986

10

## Hidden Markov Model Notation



- Underlying HMM states  $\mathbf{s} = \{s_1, \dots, s_N\}$ 
  - State at time  $t, q_t \in \mathbf{s}$
- Set of emitted observations  $\mathbf{w} = \{w_1, \dots, w_V\}$ 
  - Observation at time  $t, o_t \in \mathbf{w}$
- $\mathbf{A} = \{a_{ij}\}$ : state transition probabilities
  - $a_{ij} = P(q_{t+1} = s_j | q_t = s_i) \quad 1 \leq i, j \leq N$
- $\mathbf{B} = \{b_j(k)\}$ : observation probabilities
  - $b_j(k) = P(o_t = w_k | q_t = s_j) \equiv b_j(o_t) \quad 1 \leq j \leq N, 1 \leq k \leq V$
- $\pi = \{\pi_i\}$ : initial state distribution
  - $\pi_i = P(q_1 = s_i) \quad 1 \leq i \leq N$
- A HMM is typically written as:  $\lambda = \{\mathbf{A}, \mathbf{B}, \pi\}$



11

## POS Example with HMM Notation

- 4 states ( $s_1 = \langle \rangle, s_2 = N, s_3 = V, s_4 = \langle \rangle$ )
- $V$  word observations ( $w_{\langle \rangle}, w_{can}, w_{fish}, w_{people}, \dots$ )
- 5 input observations ( $\langle \rangle, people, can, fish, \langle \rangle$ )
  - $o_1 = w_{\langle \rangle} \quad o_2 = w_{people} \quad o_3 = w_{can} \quad o_4 = w_{fish} \quad o_5 = w_{\langle \rangle}$
- $\pi_1 = 1$



$q_1 = s_1 \quad q_2 = s_2 \quad q_3 = s_2 \quad q_4 = s_2 \quad q_5 = s_4$

$q_1 = s_1 \quad q_2 = s_3 \quad q_3 = s_3 \quad q_4 = s_3 \quad q_5 = s_4$

12

## Three Fundamental HMM Problems

---

1. Score: Given observation sequence  $\mathbf{O} = \{o_1, \dots, o_T\}$ , and HMM  $\lambda = \{\mathbf{A}, \mathbf{B}, \pi\}$ , how do we compute the probability  $P(\mathbf{O}|\lambda)$ ?
  - Forward-Backward algorithm
2. Match: Given  $\mathbf{O} = \{o_1, \dots, o_T\}$ , how do we choose the optimum underlying state sequence  $\mathbf{Q} = \{q_1, \dots, q_T\}$ ?
  - Viterbi algorithm
3. Train: How to learn ML parameter estimates for  $\lambda = \{\mathbf{A}, \mathbf{B}, \pi\}$ ?
  - Baum-Welch Estimation

13


## The Forward Algorithm

---

- Goal: compute  $P(\mathbf{O}|\lambda) = \sum_{\forall \mathbf{Q}} P(\mathbf{O}, \mathbf{Q}|\lambda)$  (brute force:  $O(TN^T)$ )
- Recursion: define *forward* variable:  $\alpha_t(i) = P(o_1, \dots, o_t, q_t = s_i|\lambda)$  (i.e., probability of seeing observations up to time  $t$ , and state  $s_i$  at time  $t$ )
  1. For  $t = 1$ ,  $\alpha_1(i) = \pi_i b_i(o_1)$   $1 \leq i \leq N$
  2. For  $t > 1$ , consider all ways of getting to current state at  $t$ 
$$\alpha_t(j) = \left[ \sum_{i=1}^N \alpha_{t-1}(i) a_{ij} \right] b_j(o_t) \quad 1 < t \leq T \quad 1 \leq j \leq N$$
  3. Finally:  $P(\mathbf{O}|\lambda) = \sum_{i=1}^N \alpha_T(i)$
- Computation is on the order of  $O(TN^2)$

14

## Forward Algorithm: $\alpha_t(i) = P(o_1, \dots, o_t, q_t = s_i | \lambda)$



15

## The Backward Algorithm

- Define **backward** variable:  $\beta_t(i) = P(o_{t+1}, \dots, o_T | q_t = s_i, \lambda)$   
(i.e., state  $s_i$  at time  $t$ , probability of seeing remaining observations)

- For  $t = T$ ,  $\beta_T(i) = 1 \quad 1 \leq i \leq N$
- For  $t < T$ , consider all ways of getting to current state at  $t$

$$\beta_t(i) = \sum_{j=1}^N a_{ij} b_j(o_{t+1}) \beta_{t+1}(j) \quad 1 \leq t < T \quad 1 \leq i \leq N$$

- Finally:  $P(\mathbf{O} | \lambda) = \sum_{i=1}^N \pi_i b_i(o_1) \beta_1(i)$


- Either forward or backward algorithm can be used to compute  $P(\mathbf{O} | \lambda)$ , but both are needed to learn model parameters

$$\alpha_t(i) \beta_t(i) = P(\mathbf{O}, q_t = s_i | \lambda)$$

16



## Backward Algorithm: $\beta_t(i) = P(o_{t+1}, \dots, o_T | q_t = s_i, \lambda)$



17

## HMM Outline

1. Score: Given observation sequence  $\mathbf{O} = \{o_1, \dots, o_T\}$ , and HMM  $\lambda = \{\mathbf{A}, \mathbf{B}, \pi\}$ , how do we compute the probability  $P(\mathbf{O}|\lambda)$ ?
  - Forward-Backward algorithm
2. Match: Given  $\mathbf{O} = \{o_1, \dots, o_T\}$ , how do we choose the optimum underlying state sequence  $\mathbf{Q} = \{q_1, \dots, q_T\}$ ?
  - Viterbi algorithm
3. Train: How to learn ML parameter estimates for  $\lambda = \{\mathbf{A}, \mathbf{B}, \pi\}$ ?
  - Baum-Welch Estimation

18

## Finding Optimal State Sequences: Viterbi Algorithm

- The Viterbi Algorithm chooses the state sequence which maximizes  $P(\mathbf{Q}|\mathbf{O}, \lambda)$  (or  $P(\mathbf{Q}, \mathbf{O}|\lambda)$ )
- Define  $\delta_t(i)$  as the highest probability along a single path to state  $s_i$  at time  $t$ , which accounts for the first  $t$  observations

$$\delta_t(i) = \max_{q_1 \dots q_{t-1}} P(q_1 \dots q_{t-1}, q_t = s_i, o_1 \dots o_t | \lambda)$$

- By induction (due to DP optimal substructure):

$$\delta_{t+1}(j) = \left[ \max_i \delta_t(i) a_{ij} \right] b_j(o_{t+1})$$

- Note similarity to the forward algorithm (except max instead of sum)
- To retrieve the best state sequence, we also keep track of the state sequence which gave the best path to state  $s_i$  at time  $t$ 
  - This is done in a separate array  $\psi_t(i)$  (i.e., pointer to best prior index)

19

## The Viterbi Algorithm

1. Initialization:  $\delta_1(i) = \pi_i b_i(o_1) \quad 1 \leq i \leq N$   
 $\psi_1(i) = 0$

2. Recursion:

$$\delta_{t+1}(j) = \max_i [\delta_t(i) a_{ij}] b_j(o_{t+1}) \quad 1 \leq t < T \quad 1 \leq j \leq N$$

$$\psi_{t+1}(j) = \arg \max_i [\delta_t(i) a_{ij}]$$

3. Termination:

$$P^* = \max_i \delta_T(i)$$


$$q_T^* = \arg \max_i \delta_T(i)$$

4. Path (state-sequence) backtracking

$$q_t^* = \psi_{t+1}(q_{t+1}^*) \quad 1 \leq t < T$$

20

## The Viterbi Algorithm



$$1) \delta_1(i) = \pi_i b_i(o_1) \quad 2) \delta_{t+1}(j) = \max_i [\delta_t(i) a_{ij}] b_j(o_{t+1}) \quad 3) P^* = \max_i \delta_T(i)$$

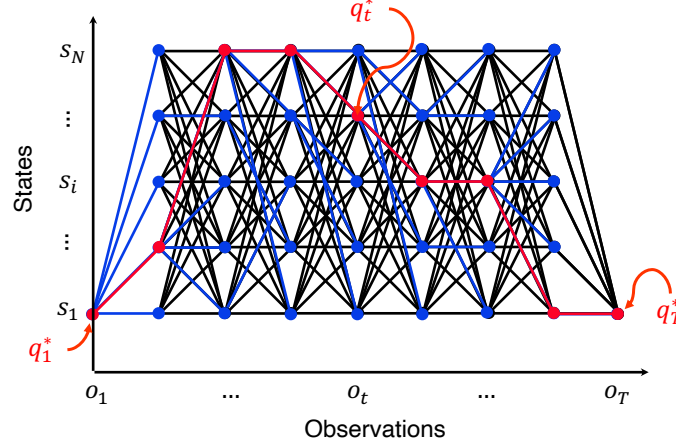

21

## The Viterbi Backtrace

- The Viterbi backtrace begins after the forward recursion completes
  - The backtrace is typically a fraction of the overall computation

$$q_T^* = \arg \max_i \delta_T(i)$$

$$q_t^* = \psi_{t+1}(q_{t+1}^*)$$



22

## HMM Outline

---

1. Score: Given observation sequence  $\mathbf{O} = \{o_1, \dots, o_T\}$ , and HMM  $\lambda = \{\mathbf{A}, \mathbf{B}, \pi\}$ , how do we compute the probability  $P(\mathbf{O}|\lambda)$ ?
  - Forward-Backward algorithm
2. Match: Given  $\mathbf{O} = \{o_1, \dots, o_T\}$ , how do we choose the optimum underlying state sequence  $\mathbf{Q} = \{q_1, \dots, q_T\}$ ?
  - Viterbi algorithm
3. Train: How to learn ML parameter estimates for  $\lambda = \{\mathbf{A}, \mathbf{B}, \pi\}$ ?
  - Baum-Welch Estimation

23

## Baum-Welch Estimation

---

- Baum-Welch estimation uses EM to determine HMM parameters
- Define  $\xi_t(i, j)$  as the probability of being in state  $s_i$  at time  $t$  and state  $s_j$  at time  $t + 1$ , given the model and observation sequence

$$\xi_t(i, j) = P(q_t = s_i, q_{t+1} = s_j | \mathbf{O}, \lambda) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{P(\mathbf{O} | \lambda)}$$

- Use  $\gamma_t(i)$  (probability of being in state  $i$  at time  $t$  given observations)

$$\gamma_t(i) = P(q_t = s_i | \mathbf{O}, \lambda) = \frac{\alpha_t(i) \beta_t(i)}{P(\mathbf{O} | \lambda)} \quad \gamma_t(i) = \sum_{j=1}^N \xi_t(i, j)$$

- Baum-Welch parameter estimates are based on expected values

$$\hat{E}(s_i \rightarrow s_j) = \sum_{t=1}^{T-1} \xi_t(i, j) \quad \hat{E}(s_j, w_k) = \sum_{\substack{t=1 \\ o_t = w_k}}^T \gamma_t(j)$$

24

## Baum-Welch Estimation Formulas

### Initialization

$$\hat{\pi}_i = \hat{E}(q_1 = s_i) = \gamma_1(i)$$

Expected number of times in state  $s_i$  at  $t = 1$

### Transition

$$\hat{a}_{ij} = \frac{\hat{E}(s_i \rightarrow s_j)}{\hat{E}(s_i \rightarrow s_*)} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}$$

Expected number of transitions from  $s_i$  to  $s_j$   
Expected number of transitions from  $s_i$


### Observation

$$\hat{b}_j(k) = \frac{\hat{E}(s_j, w_k)}{\hat{E}(s_j)} = \frac{\sum_{t=1}^T \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)}$$

Expected number of times in state  $s_j$  with symbol  $w_k$   
Expected number of times in state  $s_j$

25


## Baum-Welch State Initialization Estimation



$$\hat{\pi}_i = \hat{E}(q_1 = s_i) = \gamma_1(i) \quad \text{Expected number of times in state } s_i \text{ at } t = 1$$


26

## Baum-Welch State Transition Estimation



27

## Baum-Welch State Observation Estimation



28

## Baum-Welch Estimation Properties

- If  $\lambda = \{\mathbf{A}, \mathbf{B}, \pi\}$  is the initial model, and  $\hat{\lambda} = \{\hat{\mathbf{A}}, \hat{\mathbf{B}}, \hat{\pi}\}$  is the re-estimated model, then it can be proved that either:
  1. The initial model,  $\lambda$ , defines a critical point of the likelihood function, in which case  $\lambda = \hat{\lambda}$ , or
  2.  $P(\mathbf{O}|\hat{\lambda}) > P(\mathbf{O}|\lambda)$ : we have found a new model that is more likely to have generated the observation sequence
- Therefore, we can improve the likelihood  $P(\mathbf{O}|\lambda)$  if we iterate the re-estimation until some convergence threshold
- The resulting model is called maximum likelihood HMM
  - It is possible to over-fit parameters on a training set:  $P(\mathbf{O}|\hat{\lambda}) > P(\mathbf{O}|\lambda_{true})$

29

## Training with a Corpus

- In practice, many observation sequences  $\mathbf{O} = \{\mathbf{O}^1, \dots, \mathbf{O}^L\}$  used
  - BW estimation formulas modified to add up counts for each sequence
  - Assume that observation sequences are mutually independent

$$P(\mathbf{O}|\lambda) = \prod_{l=1}^L P(\mathbf{O}^l|\lambda)$$

- Modifications accumulate expected counts across sequences
  - Forward-backward run on each “sentence” individually, then accumulated across the entire training “corpus”

$$\hat{a}_{ij} = \frac{\hat{E}(s_i \rightarrow s_j)}{\hat{E}(s_i \rightarrow s_*)} = \frac{\sum_{l=1}^L \sum_{t=1}^{T_l-1} \xi_t^l(i, j)}{\sum_{l=1}^L \sum_{t=1}^{T_l-1} \gamma_t^l(i)} \quad \hat{b}_j(k) = \frac{\hat{E}(s_j, w_k)}{\hat{E}(s_j)} = \frac{\sum_{l=1}^L \sum_{t=1}^{T_l} \gamma_t^l(j)}{\sum_{l=1}^L \sum_{t=1}^{T_l} \gamma_t^l(j)}$$

30

## Practical HMM Issues

---

- The forward-backward and Viterbi recursions can result in long sequences of probabilities being multiplied that can cause underflow
  - *In practice, computations are performed using logprobs (e.g., your pset!)*
  - In Viterbi, multiplication of probabilities turns into sums of logprobs
  - In forward-backward both multiplications and additions are involved, so probabilities are scaled at each time frame & scale factor retained
- Often the forward-backward algorithm is not used for training, but is replaced by the simpler **Viterbi training**
  - A Viterbi “best-path” alignment is computed and then used as the basis for estimating transition probabilities, and observation parameters
  - Parameter estimation remains iterative
- For some HMM tasks, relative (and absolute) thresholds are used to eliminate unlikely hypotheses (not admissible, but practical)

31

## Final Thoughts

---

- Hidden Markov models are useful for sequential labeling tasks
  - The Viterbi algorithm is an efficient way to find the best label sequence
  - The Forward-Backward and Baum-Welch estimates enable ML training
- The mathematical formulations for HMMs were developed in 1960s
  - Applied to many disciplines with sequential data e.g., speech recognition (1970s), bioinformatics (1980s), NLP (1990s), handwriting, finance, etc.
- HMMs have been surpassed by discriminative methods such as CRFs and RNNs, but remain popular in low-resource scenarios

32



## References

---

- Readings:
  - Jurafsky & Martin, “Speech and Language Processing,” 2020 (HMMs)