

Recurrent Neural Networks

Jim Glass / MIT 6.806-6.864 / Spring 2021

1

Today's RNN Story

- RNN language models
- Training RNNs
- Gated RNNs
 - Gated recurrent units (GRUs)
 - Long short-term memories (LSTMs)
- Other architectures
 - Deep RNNs
 - Bidirectional RNNs

2

Language Modeling (Again)

- n -gram LMs represent history with the previous $n - 1$ words
 - Number of n -gram parameters increases exponentially with n

$$p(w_t | w_{t-(n-1)}, \dots, w_{t-1})$$

- An alternative model represents word history with a latent variable

$$p(w_t | w_1, \dots, w_{t-1}) \approx p(w_t | h_t)$$

- Recurrent neural networks use hidden states to capture history
 - Latent variable h_t is computed based on input x_t and h_{t-1}

$$h_t = f(x_t, h_{t-1})$$

3

A Feed-Forward LM

- Softmax output

$$\mathbf{y} = \text{softmax}(\mathbf{U}\mathbf{h} + \mathbf{c}) \quad \mathbf{U} \in \mathbb{R}^{V \times h} \quad \mathbf{c} \in \mathbb{R}^V$$

- Non-linear hidden layer

$$\mathbf{h} = f(\mathbf{W}\mathbf{v} + \mathbf{b}) \quad \mathbf{W} \in \mathbb{R}^{h \times 4d} \quad \mathbf{b} \in \mathbb{R}^h$$

- Truncated context vector

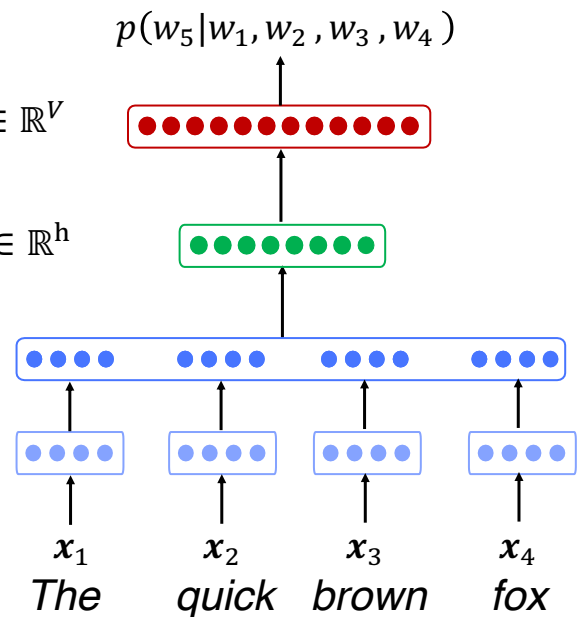
$$\mathbf{v} = [\mathbf{v}_1; \mathbf{v}_2; \mathbf{v}_3; \mathbf{v}_4] \quad \mathbf{v} \in \mathbb{R}^{4d}$$

- Word embedding vectors

$$\mathbf{v}_i = \mathbf{E}\mathbf{x}_i \quad \mathbf{v}_i \in \mathbb{R}^d \quad \mathbf{E} \in \mathbb{R}^{d \times V}$$

- One-hot input vectors

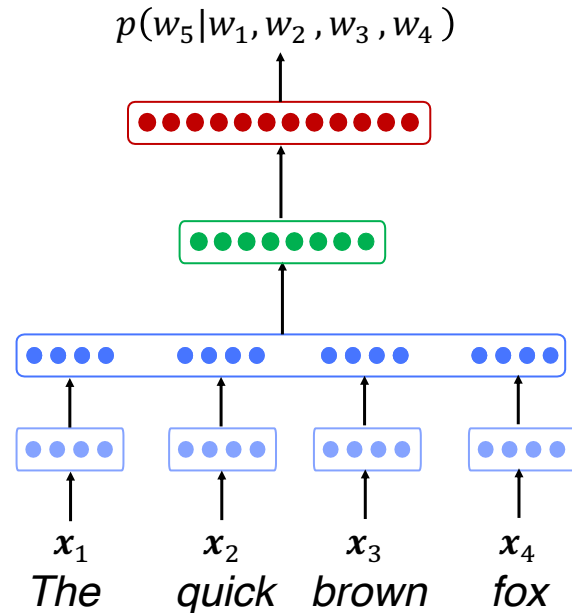
$$\mathbf{x}_i \in \{\mathbf{e}_i; 1 \leq i \leq V\}$$



4

Feed-Forward LMs

- Advantages
 - No n -gram sparsity issues
 - Better memory usage
- Disadvantages
 - Fixed context window
 - Limits ability to capture history
 - No parameter sharing in \mathbf{W}
- Need LM model to handle variable length sequential inputs




5

Recurrent Neural Networks (RNNs)

- RNNs accept variable length input sequences x_t
 - Use a hidden layer that incorporates current input and prior state
- $$\mathbf{h}_t = f(\mathbf{W}_{hh}\mathbf{h}_{t-1} + \mathbf{W}_{xh}\mathbf{x}_t + \mathbf{b}_h)$$
- Optional outputs can be produced at every step


$$\mathbf{y}_t = g(\mathbf{W}_{hy}\mathbf{h}_t + \mathbf{b}_y)$$



6

Recurrent Neural Networks (RNNs)

- Uses a hidden layer that incorporates current input and prior hidden state
- Accepts variable length sequences
- Optional outputs can be produced at every step
- Inherent parameter sharing



7

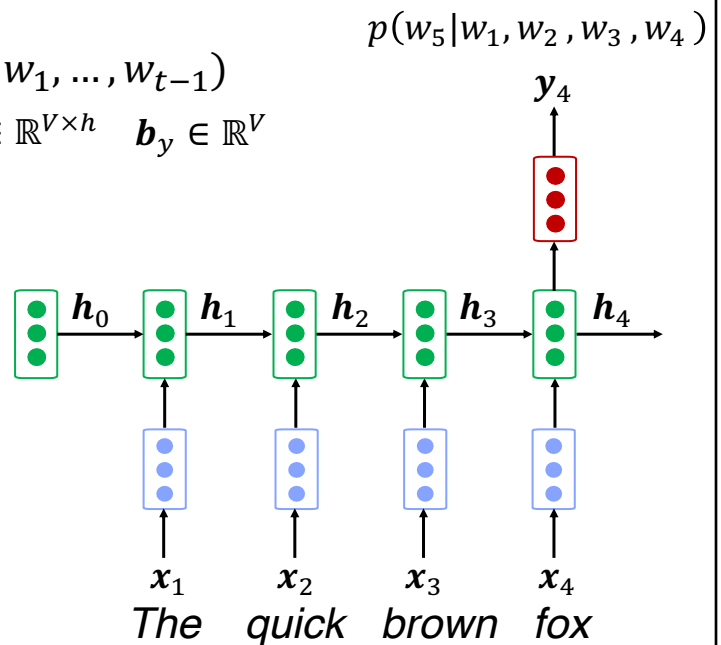
An RNN LM

- Softmax output computes $p(w_t | w_1, \dots, w_{t-1})$
- Non-linear hidden layer

$$\mathbf{h}_t = f(\mathbf{W}_{hh}\mathbf{h}_{t-1} + \mathbf{W}_{xh}\mathbf{v}_t + \mathbf{b}_h)$$

$$\mathbf{W}_{hh} \in \mathbb{R}^{h \times h} \quad \mathbf{W}_{xh} \in \mathbb{R}^{h \times d} \quad \mathbf{b}_h \in \mathbb{R}^h$$


- Word embedding vectors
 $\mathbf{v}_t = \mathbf{E}\mathbf{x}_t \quad \mathbf{v}_t \in \mathbb{R}^d \quad \mathbf{E} \in \mathbb{R}^{d \times V}$
- One-hot input vectors
 $\mathbf{x}_t \in \{\mathbf{e}_i : 1 \leq i \leq V\}$



8

RNN LMs

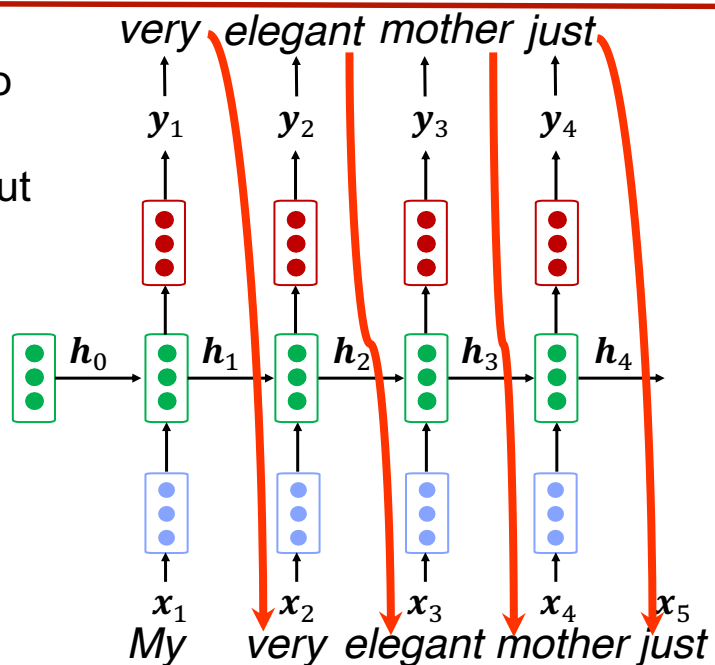
- Advantages
 - Can process variable length input
 - No truncated history
 - Model size context-independent
 - Sharing among weights
- Disadvantages
 - Recurrent computation is slow
 - Limitations on how far back it can incorporate context



9

RNN-based Language Generation


- An RNN LM can be used to *generate* text by sampling
- Sampled output is next input



10

RNN-based Sequence Labeling


- An RNN can be trained to output tags for each word
 - Part-of-speech (POS) tagging
 - Named Entity Recognition



11

RNN-based Sentence Classification


- Represent sentence as vector
- Use final hidden state as a representation of sentence
- Feed into penultimate layer for sentence classification



12

An Alternate RNN-based Sentence Classifier


- Represent sentence as vector
- Take element-wise max or mean of all hidden states
 - A simple form of *attention*
- Feed into penultimate layer for sentence classification



13

RNN-based Encoder

- Represent sentence as vector
- Take element-wise max or mean of all hidden states
 - A simple form of *attention*
- Feed into subsequent layers for downstream processing
 - Question-answering
 - Machine translation
 - Etc.




14

Sequence to Sequence Learning with Neural Networks

Ilya Sutskever
Google
ilyasu@google.com

Oriol Vinyals
Google
vinyals@google.com

Quoc V. Le
Google
qvl@google.com



NeurIPS 2014

15

Today's RNN Story

- RNN language models
- Training RNNs
- Gated RNNs
 - Gated recurrent units (GRUs)
 - Long short-term memories (LSTMs)
- Other architectures
 - Deep RNNs
 - Bidirectional RNNs

16

Training RNNs

- RNN parameters are learned on a training corpus
- Overall loss is average step-by-step cross-entropy loss

$$L(\theta) = \frac{1}{T} \sum_{t=1}^T L_t$$

- For LMs, L_t equivalent to negative log likelihood of next true word
 - Computed by taking dot-product with next one-hot vector

$$L_t = -\log p(w_t | w_1, \dots, w_{t-1}) = -\log(\mathbf{y}_i \cdot \mathbf{x}_{t+1})$$

- Parameters optimized via back-propagation and SGD


17

RNN LM Training

- The total loss is the sum of all word-by-word losses

$$L(\theta) = \frac{1}{T} \sum_{t=1}^T L_t$$

- For SGD, losses are typically accumulated in batches of sentences



18

Learning Long-Term Dependencies with Gradient Descent is Difficult

Yoshua Bengio, Patrice Simard, and Paolo Frasconi, *Student Member, IEEE*

On the difficulty of training recurrent neural networks

Razvan Pascanu

Université de Montréal, 2920, chemin de la Tour, Montréal, Québec, Canada, H3T 1J8

PASCANUR@IRO.UMONTREAL.CA

Tomas Mikolov

Speech@FIT, Brno University of Technology, Brno, Czech Republic

T.MIKOLOV@GMAIL.COM

Yoshua Bengio

Université de Montréal, 2920, chemin de la Tour, Montréal, Québec, Canada, H3T 1J8

YOSHUA.BENGIO@UMONTREAL.CA

ICML 2013

19

Training Neural Networks via Backpropagation

- Gradients propagated backwards through network to minimize loss
 - Gradients are accumulated for each parameter in a training batch
- Numerical stability issues for many layered networks
 - Exploding gradients, vanishing gradients
 - Initialization, non-linearity choices affect results
 - Techniques developed for clipping exploding gradients
 - Residual and highway connections etc. help with vanishing gradients

20



Computational Graphs

- Neural networks can be represented as a computational graph, e.g.,

$$v = Wx \quad h = f(v) \quad s = Uh$$


$$y = \text{softmax } s \quad L = -\log(y \cdot y^*)$$

- For training, each data point takes a forward and backward pass through graph
 - Gradients are accumulated for each parameter



21

Computational Graph MLP Example



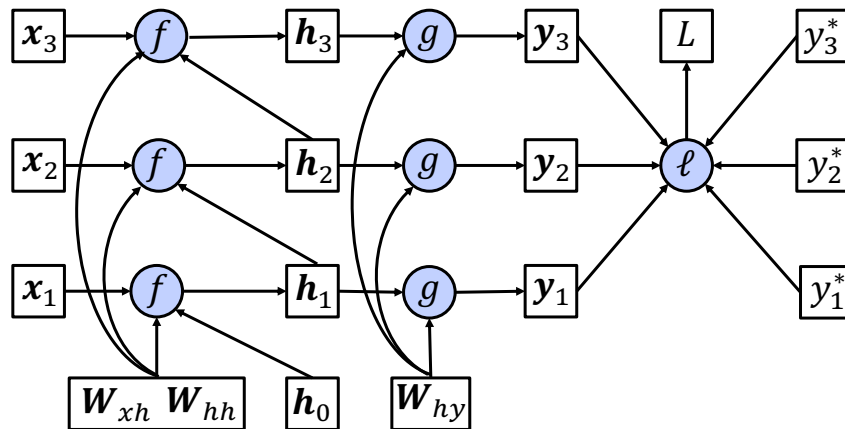
$$\frac{\partial v}{\partial W} \quad \frac{\partial h}{\partial v} \quad \frac{\partial s}{\partial h} \quad \frac{\partial s}{\partial U} \quad \frac{\partial y}{\partial s} \quad \frac{\partial L}{\partial y}$$

$$\frac{\partial L}{\partial W} = \frac{\partial L}{\partial y} \times \frac{\partial y}{\partial s} \times \frac{\partial s}{\partial h} \times \frac{\partial h}{\partial v} \times \frac{\partial v}{\partial W} \quad \frac{\partial L}{\partial U} = \frac{\partial L}{\partial y} \times \frac{\partial y}{\partial s} \times \frac{\partial s}{\partial U}$$

- Gradients are accumulated for each parameter over batch

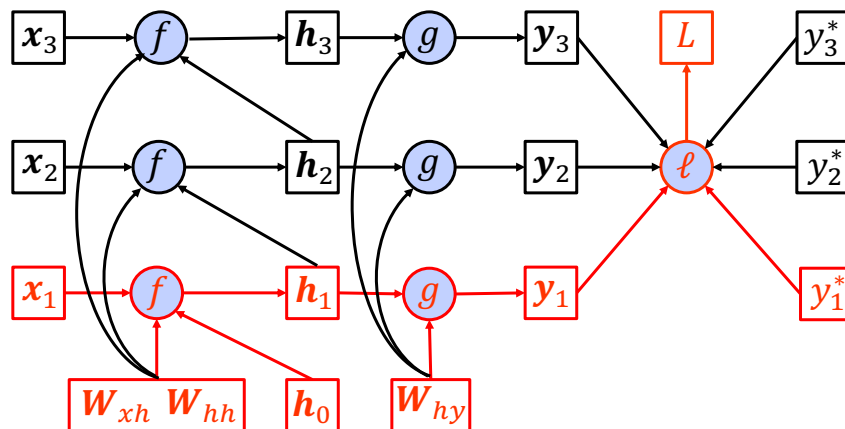
22

An RNN Computational Graph



23

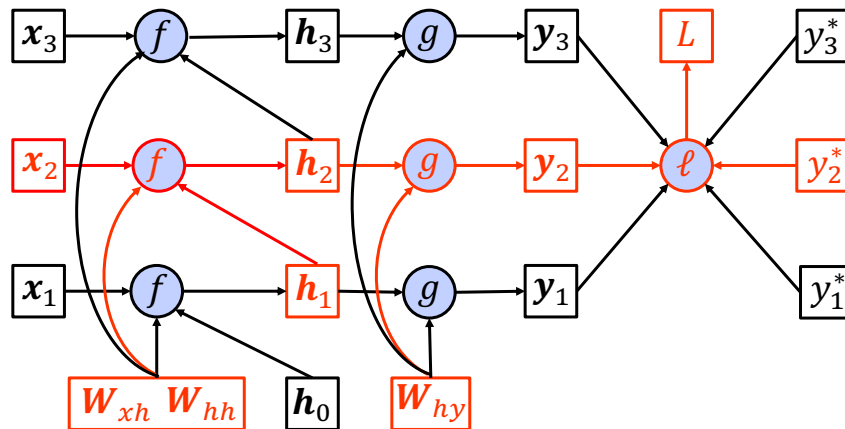
RNN Forward Pass Step 1



- Backpropagation for step 1 touches the same parameters

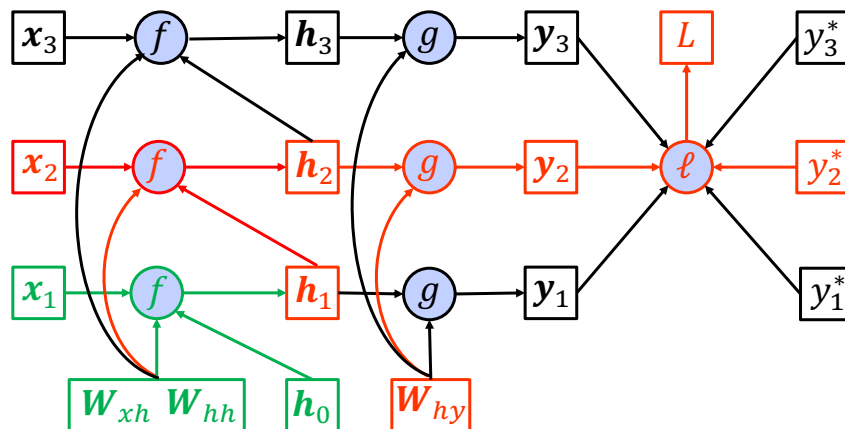
24

RNN Forward Pass Step 2



25

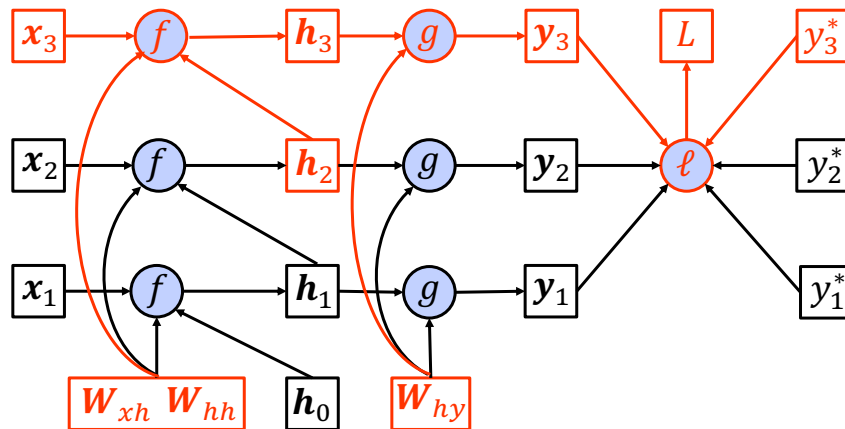
RNN Backpropagation Pass Step 2



- Backpropagation for step 2 must also consider gradients for h_1

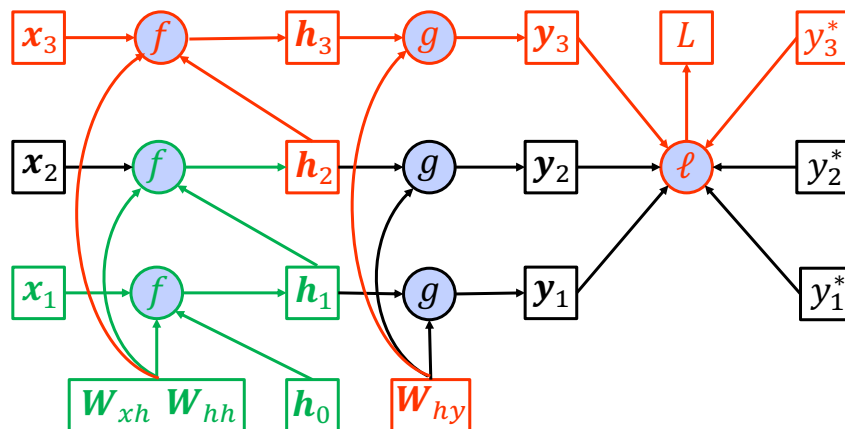
26

RNN Forward Pass Step 3



27

RNN Backpropagation Pass Step 3



- Backpropagation for step 3 must also consider partials for h_2 and h_1

28

Backpropagation Through Time

- SGD for RNNs must consider the impact of past inputs and states
 - This process is known as *Backpropagation Through Time* (BPTT)
- The gradients for longer time spans are exponential, e.g.,


$$\frac{\partial \mathbf{h}_t}{\partial \mathbf{W}_{hh}} = \sum_{i=1}^t (\mathbf{W}_{hh}^T)^{t-i} \mathbf{h}_i \quad \frac{\partial \mathbf{h}_t}{\partial \mathbf{W}_{xh}} = \sum_{i=1}^t (\mathbf{W}_{hh}^T)^{t-i} \mathbf{v}_i$$

- Potential for exploding gradients or vanishing gradients
- Since BPTT is computationally intensive for long sequences, sometimes truncated BPTT is used to save computation

29

Impact of Vanishing Gradients

- Long distance gradients are weaker and have less impact than local gradients
- Model parameters primarily learn local dependencies
- This motivated the search for RNNs that could better model long distance dependencies by some internal memory state



30

Today's RNN Story

- RNN language models
- Training RNNs
- Gated RNNs
 - Gated recurrent units (GRUs)
 - Long short-term memories (LSTMs)
- Other architectures
 - Deep RNNs
 - Bidirectional RNNs

31

Gated RNNs

- Conventional RNNs can be challenging to train
 - Long products of matrices lead to vanishing or divergent gradients
 - Effect of BPTT focuses attention on recent history
- Desiderata:
 - A *memory* mechanism to store important information over long distances
 - A *forgetting* mechanism to erase unimportant information from the model
 - A mechanism to *reset* the internal state representation
- A number of alternative RNNs attempt to address these issues
 - Gated RNNs are far more commonly used for sequence labeling
 - One of the earliest is *Long Short-Term Memory* (LSTM) RNNs
 - *Gated Recurrent Unit* (GRU) RNNs are more streamlined and faster

32

Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation



Kyunghyun Cho
Bart van Merriënboer Caglar Gulcehre
Université de Montréal
firstname.lastname@umontreal.ca

Dzmitry Bahdanau
Jacobs University, Germany
d.bahdanau@jacobs-university.de


Fethi Bougares Holger Schwenk
Université du Maine, France Université de Montréal, CIFAR Senior Fellow
firstname.lastname@lium.univ-lemans.fr find.me@on.the.web

Yoshua Bengio

Decoder



Encoder

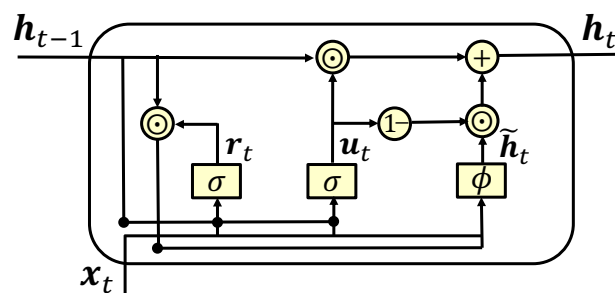


Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP), pages 1724–1734, October 25–29, 2014, Doha, Qatar. ©2014 Association for Computational Linguistics

33

Gated Recurrent Units (GRUs)

- Key distinction between regular RNNs and GRUs is gating
- Dedicated mechanisms for updating and resetting hidden state
 - **Reset** gate controls how much prior state information to remember
 - **Update** gate controls how much new state retains of old state
- Gating mechanisms are a function of current input and prior state



⊗ Element-wise multiplication

34

GRU Reset Gates


- Reset gate is used to reduce influence of \mathbf{h}_{t-1} (i.e., reset past)

$$\mathbf{r}_t = \sigma(\mathbf{W}_{hr}\mathbf{h}_{t-1} + \mathbf{W}_{xr}\mathbf{x}_t + \mathbf{b}_r) \quad \mathbf{r}_t \in \mathbb{R}^h$$

- Produces a candidate hidden state by de-weighting prior state

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W}_{hh}(\mathbf{r}_t \odot \mathbf{h}_{t-1}) + \mathbf{W}_{xh}\mathbf{x}_t + \mathbf{b}_h)$$

- For $\mathbf{r}_t \sim 1$ GRU behaves as RNN; for $\mathbf{r}_t \sim 0$ GRU behaves as MLP



35

GRU Update Gates


- Determines extent new state is old state vs new candidate state

$$\mathbf{u}_t = \sigma(\mathbf{W}_{xu}\mathbf{x}_t + \mathbf{W}_{hu}\mathbf{h}_{t-1} + \mathbf{b}_u) \quad \mathbf{u}_t \in \mathbb{R}^h$$

- Update gate \mathbf{u}_t applied in convex combination with \mathbf{h}_{t-1} and $\tilde{\mathbf{h}}_t$

$$\mathbf{h}_t = \mathbf{u}_t \odot \mathbf{h}_{t-1} + (1 - \mathbf{u}_t) \odot \tilde{\mathbf{h}}_t$$

- When $\mathbf{u}_t \sim 1$ we essentially skip time step t and remember prior state




36

LONG SHORT-TERM MEMORY

NEURAL COMPUTATION 9(8):1735-1780, 1997

Sepp Hochreiter
Fakultät für Informatik
Technische Universität München
80290 München, Germany
hochreit@informatik.tu-muenchen.de
<http://www7.informatik.tu-muenchen.de/~hochreit>

Jürgen Schmidhuber
IDSIA
Corso Elvezia 36
6900 Lugano, Switzerland
juergen@idsia.ch
<http://www.idsia.ch/~juergen>



37

Long Short-Term Memories (LSTMs)

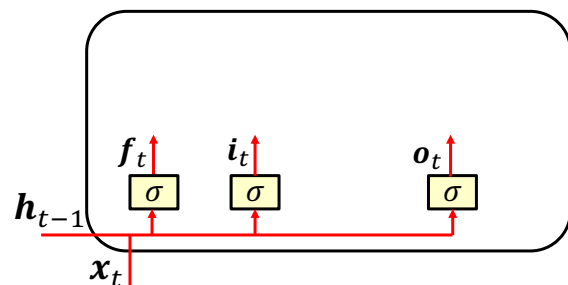
- LSTMs pre-date GRUs, but are slightly more complex (3 gates)
- Inspired by logic gates to control a *memory cell*
 - An *output* gate reads out entries from the cell
 - An *input* gate is used to read data into a cell
 - A *forget* gate is used to reset cell contents
- All three gates are a function of current input and prior state

$$i_t = \sigma(W_{hi}h_{t-1} + W_{xi}x_t + b_i)$$

$$f_t = \sigma(W_{hf}h_{t-1} + W_{xf}x_t + b_f)$$

$$o_t = \sigma(W_{ho}h_{t-1} + W_{xo}x_t + b_o)$$

$$i_t, f_t, o_t \in \mathbb{R}^h$$



38

LSTM Memory Cell


- A *candidate* memory cell is based on a regular RNN hidden state

$$\tilde{c}_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$$

- The input and forget gates are used to create the new memory cell

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

- i_t controls how much new information to take into account
- f_t controls how much old information to retain
- If $f_t \sim 1$ and $i_t \sim 0$ then old information will be retained



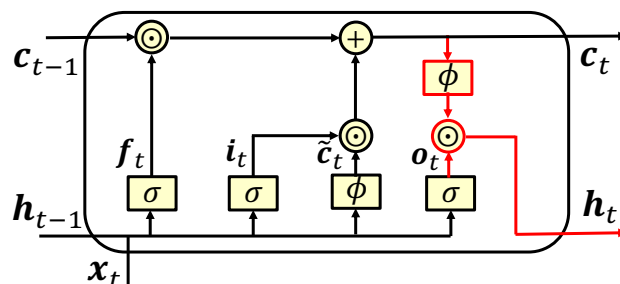
39

LSTM Hidden State

- The new hidden state is partially read from the new memory cell

$$h_t = o_t \odot \tanh(c_t)$$

- Amount retained in h_t is controlled by output gate
- \tanh ensures h_t spans interval $(-1,1)$
- If $o_t \sim 1$ pass all information through to prediction for next time step



40

GRUs vs LSTMs

- Both gated RNNs are much better able to maintain information over many timesteps compared to a vanilla RNN
- Both models have been very effective on many NLP tasks
 - LSTMs attained state-of-the-art results in the 2013-2015 time frame
 - GRUs are newer models, but have also achieved good results
- The LSTM memory cell is not bounded like the hidden state, and has demonstrated an excellent ability to count etc.
- The GRU has fewer parameters and is faster than the LSTM

41


Today's RNN Story

- RNN language models
- Training RNNs
- Gated RNNs
 - Gated recurrent units (GRUs)
 - Long short-term memories (LSTMs)
- Other architectures
 - Deep RNNs
 - Bidirectional RNNs

42

Deep RNNs

- The conventional DNN is a feed-forward network
- RNNs can be stacked in multiple layers on top of each other
 - The hidden state \mathbf{h}_t from lower RNN becomes input to the upper RNN
 - The topmost layer is responsible for generating any outputs \mathbf{y}_t
- Deep RNN layers can potentially focus on different information
 - Much more flexibility than HMMs



43

Bidirectional RNNs

- Bidirectional RNNs consist of two RNNs
 - \overrightarrow{RNN} runs in a forward direction starting from the beginning

$$\vec{\mathbf{h}}_t = \phi(\vec{\mathbf{W}}_{hh} \vec{\mathbf{h}}_{t-1} + \vec{\mathbf{W}}_{xh} \mathbf{x}_t + \vec{\mathbf{b}}_h)$$
 - \overleftarrow{RNN} runs in a backward direction starting from the end

$$\overleftarrow{\mathbf{h}}_t = \phi(\overleftarrow{\mathbf{W}}_{hh} \overleftarrow{\mathbf{h}}_{t+1} + \overleftarrow{\mathbf{W}}_{xh} \mathbf{x}_t + \overleftarrow{\mathbf{b}}_h)$$
- At time step t , the hidden state is the concatenation of $\vec{\mathbf{h}}_t$ and $\overleftarrow{\mathbf{h}}_t$


$$\mathbf{h}_t = [\vec{\mathbf{h}}_t; \overleftarrow{\mathbf{h}}_t] \quad \mathbf{h}_t \in \mathbb{R}^{2h}$$
- The bidirectional RNN output is computed like vanilla RNNs

$$\mathbf{y}_t = \phi(\mathbf{W}_{hq} \mathbf{h}_t + \mathbf{b}_q) \quad \mathbf{W}_{hq} \in \mathbb{R}^{q \times 2h}$$

44

Bidirectional RNNs

- Provides fuller context for each input token
- Bidirectional RNNs have achieved very good performance
- Requires access to entire input label sequence



45

References

- Readings:
 - Jurafsky & Martin, "Speech and Language Processing," 2020 (RNNs 9.2-9.3)

46